

CS110 Midterm Examination

This is a closed book, closed note, closed computer exam (although you can use your single double-sided cheat sheet). You have 80 minutes to complete all problems. You don't need to **#include** any header files, and you needn't guard against any errors unless specifically instructed to do so. Understand that most points are awarded for concepts taught in CS110. If you're taking the exam remotely, you can call me at 415-205-2242 should you have any questions.

Good luck!

SUNet ID (username): _____@**stanford.edu**

Last Name: _____

First Name: _____

I accept the letter and spirit of the honor code.

[signed] _____

	Score	Grader
1. findexec	[16] _____	_____
2. Short Answer Questions	[14] _____	_____
Total	[30] _____	_____

Relevant Prototypes

```
// filesystem access
int open(const char *path, int oflag, ...); // returns descriptor
ssize_t read(int fd, char buffer[], size_t len); // returns num read, 0 at eof
ssize_t write(int fd, char buffer[], size_t len); // returns num written
int close(int fd); // ignore retval
int pipe(int fds[]); // argument should be array of length 2, ignore retval
int dup2(int old, int new); // ignore retval
#define STDIN_FILENO 0
#define STDOUT_FILENO 1
#define STDERR_FILENO 2

// exceptional control flow and multiprocessing
pid_t fork();
pid_t waitpid(pid_t pid, int *status, int flags);
typedef void (*sighandler_t)(int sig);
sighandler_t signal(int signum, sighandler_t handler); // ignore retval
int execvp(const char *path, char *argv[]); // ignore retval
int kill(pid_t pid, int signal); // ignore retval
#define WIFEXITED(status) // macro
#define WEXITSTATUS(status) // macro
```

Problem 1: findexec [16 points]

Implement the **findexec** function, which has the following prototype:

```
void findexec(char *root, char *pattern, char *command[]);
```

This **findexec** function runs two sister processes—the first running **find** and a second running **xargs**—such that the standard output of the first leads to the standard input of the second. The **root** and **pattern** parameters shape how the first process does its job, and **command** is a **NULL**-terminated vector of arguments that shape how the second process does its job. Restated, the **root**, **pattern**, and **command** parameters prompt **findexec** to do the same thing that the following command would have achieved in **stsh**:

```
stsh> find root -name pattern -print | xargs command
```

Your implementation must do the following:

- construct the argument vector for the first process to include the supplied **root** and **pattern** variables
- construct the argument vector for the second process where "**xargs**" has been prepended to the series of tokens residing within **command**
- create two new sister processes such that the stdout output of the first is wired to the standard input of the second
- transform the first to execute **find**, and the second to execute **xargs**
- close all unused file descriptors
- wait for each of the two processes to finish before returning

You may assume all system calls work as intended and need not do any error checking at all. You may not use **pipeline** or **subprocess** from lecture or the assignment set. In fact, I basically want you to re-implement something akin to Assignment 3's **pipeline**, but specific to **find** and **xargs**. Note that we're not asking you to implement **find** or **xargs**, as we've already implemented our own versions of them in lecture and in discussion section.

Use the next two pages to present your implementation.

Problem 1: findexec [continued]

```
size_t countTokens(const char *command[]) {  
    size_t count;  
    for (count = 0; command[count] != NULL; count++) {}  
    return count;  
}  
  
void findexec(char *root, char *pattern, char *command[]) {
```

Problem 1: findexec [continued]

Problem 2: Short Answer Questions [14 points]

Unless otherwise noted, your answers to the following questions should be 75 words or fewer. You needn't write in complete sentences provided it's clear what you're saying. Full credit will only be given to the best of responses. Just because everything you write is true doesn't mean you get all the points.

- a. [2 points] **remove** is a C library function that removes a name from the file system. If the supplied name was the last one to identify the file, then the file itself is truly deleted and its resources donated back for reuse. In the context of your **assign2** filesystem design, explain how the file system would need to be updated to fully realize a call to **remove**.

- b. [2 points] Recall that the **struct inode** from **assign2** looked like this:

```
struct inode { // some fields irrelevant to the problem are omitted
    uint16_t i_mode;      // bit vector of file type and permissions
    uint8_t  i_size0;     // most significant byte of size
    uint16_t i_size1;     // lower two bytes of size
    uint16_t i_addr[8];   // device addresses constituting file
};
```

One CS110 student once proposed the following idea: for files with sizes that are just slightly larger than a perfect multiple of the block size (e.g. 1027, when the block size is 256), those last few bytes (e.g. the last three bytes of a 1027-byte file) could be stored in the inode itself. Describe how you would support this optimization so that entire blocks needn't be allocated just to store a few bytes of payload.

- c. [2 points] Referring to your implementation of **findexec** in Problem 1, identify the one **close** call that, if omitted, would prevent **findexec** from doing its job. Then explain why that one **close** call is so crucial to everything working as expected.
- d. [2 points] Your implementations of **farm** and **stsh** each relied on signal handlers—**farm** relied on a **SIGCHLD** handler to identify stopped processes, and **stsh** relied on a **SIGCHLD** handler to identify processes that have stopped, exited, crashed, or exited normally. In fact, **farm** could have been implemented just as easily without custom handlers, whereas **stsh** really needed them. Explain why this is true.

- e. [2 points] When establishing a new process group for a pipeline of two or more commands (as with `echo "abcdefgh" | ./conduit --count 4`), your `stsh` implementation needed to call `setpgid` in both the parent and in each of the children ("in order to avoid some race conditions", as the handout stated it). Describe the race condition that could cause problems if the first child didn't call `setpgid` and instead just relied on the parent to call it before moving on to the create the second child.
- f. [2 points] Explain what the scheduler does when a program makes an otherwise valid call to `read` at a time when no data is available. Further explain what the scheduler does so that it's informed when data does become available.
- g. [2 points] Many students asked if one signal handler can be interrupted by a signal of a different type. Describe a simple coding experiment you could run to list all the signals capable of interrupting a `SIGCHLD` handler.