

Lab Solution 2: Multiprocessing and Unix Tools

The first three problems are questions that could easily appear on a midterm or final exam. The exercises beyond that are designed to help you master the development tools even more so than you have already.

Students are encouraged to share their ideas on Piazza, which has been set up [right here](#). SCPD students are welcome to reach out to me directly if they have questions that can't be properly addressed without being physically present for a discussion section.

Before starting, go ahead and clone the `lab2` folder, which contains a working solution to Problem 1 and fodder for some of the later problems that have you experiment with some new development tools.

```
poohbear@myth15:~$ hg clone /usr/class/cs110/repos/lab2/shared lab2
poohbear@myth15:~$ cd lab2
poohbear@myth15:~$ make
```

Solution 1: Implementing `exargs`

`exargs` (designed here to emulate the `xargs` builtin—type `man xargs` for the full read) reads tokens from standard input (delimited by spaces and newlines), and executes the command with any initial arguments followed by the tokens read from standard input. `exargs` (or `xargs`) is useful when one program is needed to programmatically generate the argument vector for a second program. (Understand that the builtin `xargs` is much more sophisticated than our `exargs`; `xargs` has all kinds of parsing options and flags, but our `exargs` has no such bells and whistles.)

To illustrate the basic idea, consider the `factor` program, which prints out the prime factorizations of all of its numeric arguments, as with:

```
poohbear@myth15:~$ factor 720
720: 2 2 2 2 3 3 5
poohbear@myth15:~$ factor 9 16 2047 870037764750
9: 3 3
16: 2 2 2 2
2047: 23 89
870037764750: 2 3 3 5 5 5 7 7 7 11 11 11 11 11
```

```
poohbear@myth15:~$
```

To see how `exargs` works, check this out:

```
poohbear@myth15:~$ printf "720" | ./exargs factor
720: 2 2 2 2 3 3 5
poohbear@myth15:~$ printf "2047 1000\n870037764750" | ./exargs factor 9 16
9: 3 3
16: 2 2 2 2
2047: 23 89
1000: 2 2 2 5 5 5
870037764750: 2 3 3 5 5 5 7 7 7 7 11 11 11 11 11
poohbear@myth15:~$
```

Note that the first process in the pipeline—the `printf`—is a brute force representative of an executable capable of supplying or extending the argument vector of a second executable—in this case, `factor`—through `exargs`. Of course, the two executables needn't be `printf` or `factor`; they can be anything that works. If, for example, I'm interested in exposing just how much code I had to write for my own `assign2` solution (see if you can access the folder :)), I might use `exargs` to do this:

```
poohbear@myth30:~$ ls /usr/class/cs110/staff/master_repos/assign2/*.c | ./exargs
wc --chars --lines --max-line-length
 78 1792   90 /usr/class/cs110/staff/master_repos/assign2/chksumfile.c
 35 1178  121 /usr/class/cs110/staff/master_repos/assign2/directory.c
266 8015  111 /usr/class/cs110/staff/master_repos/assign2/diskimageaccess.c
 31  731   86 /usr/class/cs110/staff/master_repos/assign2/disking.c
 35 1193  144 /usr/class/cs110/staff/master_repos/assign2/file.c
 72 2751  134 /usr/class/cs110/staff/master_repos/assign2/inode.c
 33  987  152 /usr/class/cs110/staff/master_repos/assign2/pathname.c
 45 1287   91 /usr/class/cs110/staff/master_repos/assign2/unixfilesystem.c
595 17934 152 total
```

As a whiteboard exercise, construct the entire implementation of the `exargs` program, relying on the following utility function to parse all of standard input around newlines and spaces:

```
static void pullAllTokens(istream& in, vector<string>& tokens) {
    while (true) {
```

```

    string line;
    getline(in, line);
    if (in.fail()) break;
    istringstream iss(line);
    while (true) {
        string token;
        getline(iss, token, ' ');
        if (iss.fail()) break;
        tokens.push_back(token);
    }
}
}
}

```

You need not perform any error checking on user input, and you can assume that all system calls succeed. Implement the entire program to return 0 if and only if the command executed by `exargs` exits normally with status code 0, and return 1 otherwise.

```

int main(int argc char *argv[]) {

```

*My solution is inlined below, and it was also included in the lab2 repo you need to refer to for Problems 4 and 5 below. The `fork`, `execvp`, and **`waitpid`** contributions to the solution are of paramount importance. For more information about `memcpy` and `transform`, click [here](#) and [here](#).*

```

int main(int argc, char *argv[]) {
    vector<string> tokens;
    pullAllTokens(cin, tokens);
    pid_t pid = fork();
    if (pid == 0) {
        char *exargsv[argc + tokens.size()];
        memcpy(exargsv, argv + 1, (argc - 1) * sizeof(char *));
        transform(tokens.cbegin(), tokens.cend(), exargsv + argc - 1,
            [](const string& str) { return const_cast<char *>(str.c_str());
    });
    exargsv[argc + tokens.size() - 1] = NULL;
    execvp(exargsv[0], exargsv);
    cerr << exargsv[0] << ": command not found" << endl;

```

```

        exit(0);
    }

    int status;
    waitpid(pid, &status, 0);

    return status == 0 ? 0 : 1; // trivia: if all of status is 0, then child
    exited normally with code 0
}

```

Solution 2: Incorrect Output File Redirection

The `publish` user program takes an arbitrary number of filenames as arguments and attempts to publish the date and time (via the `date` executable that ships with all versions of Unix and Linux). `publish-to-all` is built from the following source:

```

static void publish(const char *name) {
    printf("Publishing date and time to file named \"%s\".\n", name);
    int outfile = open(name, O_WRONLY | O_CREAT | O_TRUNC, 0644);
    dup2(outfile, STDOUT_FILENO);
    close(outfile);
    if (fork() > 0) return;
    char *argv[] = { "date", NULL };
    execvp(argv[0], argv);
}

int main(int argc, char *argv[]) {
    for (size_t i = 1; i < argc; i++) publish(argv[i]);
    return 0;
}

```

Someone with a fractured understanding of processes, descriptors, and file redirection might expect the program to have printed something like this:

```

myth4> ./publish one two three four
Publishing date and time to file named "one".
Publishing date and time to file named "two".

```

```
Publishing date and time to file named "three".
```

```
Publishing date and time to file named "four".
```

However, that's not what happens. Questions:

- What text is actually printed?
 - *Here's what happens when I launch the above executable from myth30:*

```
poohbear@myth30:~$ publish-to-all one two three four
```

```
Publishing date and time to file named "one".
```

- What do each of the four files contain?
 - *Contents of one (similar for two and three, lines can be interchanged):*

```
Publishing date and time to file named "two".
```

```
Wed Apr 19 10:14:14 PDT 2017
```

- *Contents of four:*

```
Wed Apr 19 10:14:14 PDT 2017
```

- How should the program be rewritten so that it works as intended?
 - *Because the child processes (and only the child processes) should be redirecting, you should open, dup2, and close in child-specific code. A happy side effect of the change is that you never muck with STDOUT_FILENO in the parent if you confine the redirection code to the child.*

```
static void publish(const char *name) {
    printf("Publishing date and time to file named \"%s\".\n", name);
    if (fork() > 0) return;
    int outfile = open(name, O_WRONLY | O_CREAT | O_TRUNC, 0644);
    dup2(outfile, STDOUT_FILENO);
    close(outfile);
    char *argv[] = { "date", NULL };
}
```

```
execvp(argv[0], argv);  
}
```

Solution 3: Short Answer Questions

Here are a bunch of short answer questions that have appeared on past CS110 midterms and final exams.

- Recall that one vnode table and one file entry table are maintained on behalf of all processes, but that each process maintains its own file descriptor table. What problems would result if just one file descriptor table were maintained on behalf of all processes?
 - *If all processes referred to the same descriptor table, one process—perhaps one owned by another user—could close or otherwise manipulate the file session behind the back of the true owner. That would certainly lead to unintended consequences, and it would also lead to security concerns (e.g. a child process might inherit access to a session leading to a file with private information).*
- Your terminal can be configured so that a process **dumps core**—that is, generates a data file named `core`—whenever it crashes (because it seg faults, for instance.) This `core` file can be loaded into and analyzed within `gdb` to help identify where and why the program is crashing. Assuming we can modify the program source code and recompile, how might you **programmatically** generate a core dump at specific point in the program while allowing the process to continue executing? (Your answer might include a **very, very short** code snippet to make its point.)
 - *Call `fork` at the line of interest, and have the child intentionally segfault (e.g. `*(int *) NULL = 14;` or `raise(SIGSEGV);` on the line immediately after the `fork`.*
- The `fork` system call creates a new process with an independent address space, where the contents of the parent's address space are replicated—in a sense, `memcpy`'ed—into the address space of the clone. If, however, a **copy-on-write** implementation strategy is adopted, then both parent and child share the same address space and only start to piecemeal split into two independent address spaces as one process makes changes that shouldn't be reflected in the other. In general, most operating systems adopt a copy-on-write approach, even though it's more difficult to implement. Defend that approach given what we've learned in class so far.
 - *The vast majority of `fork` calls lead the child process to an `execvp` call, where the child's address space is gutted 100% and replaced with a brand new one. `fork`'s defense for the **copy-of-write** model would be that it's a lot of work to replicate an entire address space only for it to be discarded a few lines later when `execvp` is called.*

Solution 4: Experimenting with `strace`

`strace` is an advanced development tool that programmers use to determine what system calls contribute to the execution of a program (generally because the program is malfunctioning, and they're curious if any failing system calls are the root cause of the problem). If,

for instance, you want to see how sleep 10 works behind the scenes, you gain a lot by typing this in and following along:

```
// lots of calls omitted in the name of brevity

fstat(3, {st_mode=S_IFREG|0644, st_size=4431376, ...}) = 0
mmap(NULL, 4431376, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f62a18ae000
close(3)                                = 0
nanosleep({10, 0}, NULL)                = 0
close(1)                                = 0
close(2)                                = 0
exit_group(0)                           = ?
+++ exited with 0 +++

poohbear@myth30:~$ poohbear@myth30:~$ strace sleep 10
execve("/bin/sleep", ["sleep", "10"], [/ * 27 vars */]) = 0
brk(0)                                   = 0x1636000
access("/etc/ld.so.nohwcap", F_OK)       = -1 ENOENT (No such file or directory)
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x7f62a22d0000
access("/etc/ld.so.preload", R_OK)      = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3

// lots of calls omitted in the name of brevity

fstat(3, {st_mode=S_IFREG|0644, st_size=4431376, ...}) = 0
mmap(NULL, 4431376, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f62a18ae000
close(3)                                = 0
nanosleep({10, 0}, NULL)                = 0
close(1)                                = 0
close(2)                                = 0
exit_group(0)                           = ?
+++ exited with 0 +++

poohbear@myth30:~$
```

A typical `strace` run starts with an `execve` (which my shell uses instead of `execvp`), and then

works through all of these systemsy things to load C++ libraries, build the heap segment, etc., until it reaches the crucial `nanosleep` call, which is the call that halts the process for 10 seconds. You see gestures to system calls that have come up in CS107, CS110, and your first two assignments: `execve`, `access`, `mmap`, `open`, and `close`.

If you're interested only in a particular subset of system calls, you can identify those of interest when invoking `strace` using the `-e trace=<csv-list-of-syscall-name>`, as with this:

```
// lots of calls omitted in the name of brevity

munmap(0x7f2ec1242000, 154308)      = 0
mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0)
= 0x7f2ec1267000
munmap(0x7f2ec1267000, 4096)      = 0
mmap(NULL, 4431376, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f2ebfdd4000
mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0)
= 0x7f2ec1267000
ARCHIVE.README  cgi-bin include lecture-examples lib local  private_data
README.md repos  samples staff  tools WWW
munmap(0x7f2ec1267000, 4096)      = 0
+++ exited with 0 +++

poohbear@myth30:~$poohbear@myth30:~$ strace -e trace=read ls /usr/class/cs110
read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\0[\0\0\0\0\0\0"... , 832)
= 832
read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\0`\34\0\0\0\0\0\0"... ,
832) = 832
read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P \2\0\0\0\0\0"... , 832)
= 832
read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\260\27\0\0\0\0\0\0"... ,
832) = 832
read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\320\16\0\0\0\0\0\0"... ,
832) = 832
read(3, "\177ELF\2\1\1\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\300\20\0\0\0\0\0\0"... ,
832) = 832
read(3, "nodev\tsysfs\nnodev\trootfs\nnodev\ttr"... , 1024) = 344
read(3, "", 1024)                  = 0
```



```

ARCHIVE.README    cgi-bin  include  lecture-examples  lib  local    private_data
README.md  repos    samples  staff    tools  WWW

+++ exited with 0 +++

poohbear@myth30:~$ strace -e trace=mmap,munmap ls /usr/class/cs110

mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x7f2ec1268000

mmap(NULL, 154308, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f2ec1242000

mmap(NULL, 2238192, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) =
0x7f2ec0e25000

mmap(0x7f2ec1044000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE,
3, 0x1f000) = 0x7f2ec1044000

mmap(0x7f2ec1046000, 5872, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS,
-1, 0) = 0x7f2ec1046000

mmap(NULL, 2126336, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) =
0x7f2ec0c1d000

mmap(0x7f2ec0e23000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE,
3, 0x6000) = 0x7f2ec0e23000

// lots of calls omitted in the name of brevity

munmap(0x7f2ec1242000, 154308)          = 0

mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x7f2ec1267000

munmap(0x7f2ec1267000, 4096)          = 0

mmap(NULL, 4431376, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f2ebfdd4000

mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x7f2ec1267000

ARCHIVE.README    cgi-bin  include  lecture-examples  lib  local    private_data
README.md  repos    samples  staff    tools  WWW

munmap(0x7f2ec1267000, 4096)          = 0

+++ exited with 0 +++

poohbear@myth30:~$

```

Take some time to experiment with `strace`, as it'll help you with your Assignment 3 work (because you implement a light version of `strace` as part of it-surprise!). Type in each of these commands from any directory and see what you get:

- `strace date`
- `strace -e trace=write date`

- `strace -e trace=clone,execve date`

Note that calls to `fork` are reframed as calls to `clone`, and calls to `execvp` are reframed as calls to `execve`. Here's what I get when I type in the three commands above:

```
// many calls omitted for brevity

write(1, "Thu Apr 20 12:11:25 PDT 2017\n", 29Thu Apr 20 12:11:25 PDT 2017
) = 29
close(1)                                = 0
munmap(0x7fc79f4c2000, 4096)             = 0
close(2)                                = 0
exit_group(0)                           = ?
+++ exited with 0 +++

poohbear@myth15:~$ strace -e trace=write date
write(1, "Thu Apr 20 12:14:59 PDT 2017\n", 29Thu Apr 20 12:14:59 PDT 2017
) = 29
+++ exited with 0 +++

poohbear@myth15:~$ strace -e trace=clone,execve date
execve("/bin/date", ["date"], [/ * 27 vars */]) = 0
Thu Apr 20 12:15:45 PDT 2017
+++ exited with 0 +++

poohbear@myth15:~$ poohbear@myth15:~$ strace date
execve("/bin/date", ["date"], [/ * 27 vars */]) = 0
brk(0)                                    = 0x1a3b000
access("/etc/ld.so.nohwcap", F_OK)       = -1 ENOENT (No such file or directory)
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) =
0x7fc79f4c3000
access("/etc/ld.so.preload", R_OK)       = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=172263, ...}) = 0
mmap(NULL, 172263, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7fc79f498000

// many calls omitted for brevity
```

```

write(1, "Thu Apr 20 12:11:25 PDT 2017\n", 29Thu Apr 20 12:11:25 PDT 2017
) = 29
close(1)                                = 0
munmap(0x7fc79f4c2000, 4096)            = 0
close(2)                                = 0
exit_group(0)                           = ?
+++ exited with 0 +++
poohbear@myth15:~$ strace -e trace=write date
write(1, "Thu Apr 20 12:14:59 PDT 2017\n", 29Thu Apr 20 12:14:59 PDT 2017
) = 29
+++ exited with 0 +++
poohbear@myth15:~$ strace -e trace=clone,execve date
execve("/bin/date", ["date"], [/* 27 vars */]) = 0
Thu Apr 20 12:15:45 PDT 2017
+++ exited with 0 +++
poohbear@myth15:~$

```

Rather than list out specific system calls via `-e trace=<csv-list-of-syscall-name>`, `strace` allows you to specify a family of system calls, as per `strace`'s man page, e.g. `-e trace=file`, or `-e trace=memory`.

Now type these commands from any directory and see what you get:

- `strace -e trace=file ls /usr/class/cs110`
- `strace -e trace=memory ls /usr/class/cs110`
- `strace -e trace=desc ls /usr/class/cs110`

Rather than list the output, I'll just be clear that the first command (with `trace=file`) should confine itself to `execve`, `access`, `open`, `statfs`, `stat`, and `openat`. Every single one of these system calls has one or more arguments that is the name of some file (e.g. `"ls"`, `"/usr/class/cs110"`, `"/proc/filesystems"`). Note that `read`, `write`, and `close` don't show up, since those are descriptor-related, not file-related.

Including `trace=memory` lists all operations that introduce and configured new memory segments: `brk`, `mmap`, `mprotect`, and `munmap`. And `trace=desc` lists anything that returns a descriptor and/or accepts one or more as arguments: `mmap` (again), `open`, `fstat`, `close`, `read`, `openat`, `getdents`, and `write`.

See what happens when you try to launch something that can't be launched, because the

specified file isn't an executable:

- `strace /usr/class/cs110/WWW`
- `strace /usr/class/cs110/WWW/index.html`

Look at this, and see how `execvp` fails! And note that the error messages are printed to file descriptor 2 (e.g. `STDERR_FILENO`).

```
poohbear@myth15:~$ strace /usr/class/cs110/WWW
execve("/usr/class/cs110/WWW", ["/usr/class/cs110/WWW"], [/* 27 vars */]) = -1
EACCES (Permission denied)
write(2, "strace: exec: Permission denied\n", 32strace: exec: Permission denied
) = 32
exit_group(1)                                = ?
+++ exited with 1 +++
poohbear@myth15:~$ strace /usr/class/cs110/WWW/index.html
execve("/usr/class/cs110/WWW/index.html", ["/usr/class/cs110/WWW/index.html"],
[/* 27 vars */]) = -1 EACCES (Permission denied)
write(2, "strace: exec: Permission denied\n", 32strace: exec: Permission denied
) = 32
exit_group(1)                                = ?
+++ exited with 1 +++
poohbear@myth15:~$
```

Just for fun, try `strace strace`, with:

- `strace strace ls`

This lists all of the system calls called by `strace`, but not those associated with the grandchild `ls`-running process.

Finally, descend into your `lab2` folder and trace a few things there:

- `printf "4 6 8" | strace -e trace=clone,execve ./exargs factor`
- `printf "4 6 8" | strace -f -e trace=clone,execve ./exargs factor`

Scan `strace`'s man page for information about the `-f` flag, and be sure you understand why the output of the second command includes so many calls to `execve`.

Here's the output of both:

```

poohbear@myth15:~/lab2$ printf "4 6 8" | strace -e trace=clone,execve ./exargs
factor
execve("./exargs", [ "./exargs", "factor" ], [ /* 27 vars */ ]) = 0
clone(child_stack=0, flags=CLONE_CHILD_CLEARTID|CLONE_CHILD_SETTID|SIGCHLD,
child_tidptr=0x7f8460c8aa50) = 11290
4: 2 2
6: 2 3
8: 2 2 2
--- SIGCHLD {si_signo=SIGCHLD, si_code=CLD_EXITED, si_pid=11290, si_status=0,
si_utime=0, si_stime=0} ---
+++ exited with 0 +++
poohbear@myth15:~/lab2$ printf "4 6 8" | strace -f -e trace=clone,execve ./
exargs factor
execve("./exargs", [ "./exargs", "factor" ], [ /* 27 vars */ ]) = 0
clone(Process 11354 attached
child_stack=0, flags=CLONE_CHILD_CLEARTID|CLONE_CHILD_SETTID|SIGCHLD,
child_tidptr=0x7f8f3af35a50) = 11354
[pid 11354] execve("/afs/ir/users/p/o/poohbear/bin/factor", ["factor", "4", "6",
"8"], [ /* 27 vars */ ]) = -1 ENOENT (No such file or directory)
[pid 11354] execve("/afs/ir/users/p/o/poohbear/local/bin/factor", ["factor",
"4", "6", "8"], [ /* 27 vars */ ]) = -1 ENOENT (No such file or directory)
[pid 11354] execve("/afs/ir/users/p/o/poohbear/.cabal/bin/factor", ["factor",
"4", "6", "8"], [ /* 27 vars */ ]) = -1 ENOENT (No such file or directory)
[pid 11354] execve("/usr/class/cs110/tools/factor", ["factor", "4", "6", "8"],
[ /* 27 vars */ ]) = -1 ENOENT (No such file or directory)
[pid 11354] execve("/usr/class/cs110/staff/bin/factor", ["factor", "4", "6",
"8"], [ /* 27 vars */ ]) = -1 ENOENT (No such file or directory)
[pid 11354] execve("/usr/local/sbin/factor", ["factor", "4", "6", "8"], [ /* 27
vars */ ]) = -1 ENOENT (No such file or directory)
[pid 11354] execve("/usr/local/bin/factor", ["factor", "4", "6", "8"], [ /* 27
vars */ ]) = -1 ENOENT (No such file or directory)
[pid 11354] execve("/usr/sbin/factor", ["factor", "4", "6", "8"], [ /* 27 vars
*/ ]) = -1 ENOENT (No such file or directory)
[pid 11354] execve("/usr/bin/factor", ["factor", "4", "6", "8"], [ /* 27 vars */
]) = 0
4: 2 2

```

```

6: 2 3
8: 2 2 2

[pid 11354] +++ exited with 0 +++
--- SIGCHLD {si_signo=SIGCHLD, si_code=CLD_EXITED, si_pid=11354, si_status=0,
si_utime=0, si_stime=0} ---
+++ exited with 0 +++
poohbear@myth15:~/lab2$

```

You should note a few things:

- *strace* doesn't read from standard input, so it doesn't compete for the material produced by the *printf*. That's good, because that material is intended for *exargs*!
- Only when you include *-f* do you see the additional *execve* calls. *strace* is also clear when the system calls are executing in a process other than the primary; it prefaces the call with something like *[pid 11354]*.
- The implementation of *execvp* (which our *exargs* relies on) loops over **every entry** in your path, appends factor to it, and see if *execve* works out. That's why you see all but the last call in the stream of *execve* calls returning -1.

Solution 5: Using valgrind and gdb with fork

You know by this point that my own solution to Problem 1 resides within that *lab2* folder you cloned. We're going to use this problem as a vehicle for learning how to use *gdb* to step through processes that split into multiple ones.

This exercise is framed in terms of an intentionally buggy version of *exargs.cc*, which I've placed in *buggy-exargs.cc*. I've done bad things in *buggy-exargs.cc* to make sure that *buggy-exargs* fails miserably. Here's a diff between the buggy version and the correct one:

```

poohbear@myth30:$ diff buggy-exargs.cc exargs.cc
47,48c47,48
<     char **exargsv = NULL;
<     memcpy(exargsv, argv + 1, (argc - 1) * sizeof(char *));
---
>     char *exargsv[argc + tokens.size()];
>     memcpy(exargsv, argv + 1, (argc - 1) * sizeof(char *));
poohbear@myth30:~/lab2$

```

Run the following commands to reaffirm that the correct version runs clean under *valgrind* (aside from the suppressed errors I spoke of in the Assignment 1 handout):

```

poohbear@myth30:~/lab2$ printf "1 2 3 4" | ./exargs factor
1:
2: 2
3: 3
4: 2 2

poohbear@myth30:~/lab2$ printf "1 2 3 4" | valgrind ./exargs factor
==12801== Memcheck, a memory error detector
==12801== Copyright (C) 2002-2013, and GNU GPL'd, by Julian Seward et al.
==12801== Using Valgrind-3.10.1 and LibVEX; rerun with -h for copyright info
==12801== Command: ./exargs factor
==12801==
1:
2: 2
3: 3
4: 2 2
==12801==
==12801== HEAP SUMMARY:
==12801==      in use at exit: 72,704 bytes in 1 blocks
==12801==    total heap usage: 13 allocs, 12 frees, 73,011 bytes allocated
==12801==
==12801== LEAK SUMMARY:
==12801==      definitely lost: 0 bytes in 0 blocks
==12801==      indirectly lost: 0 bytes in 0 blocks
==12801==      possibly lost: 0 bytes in 0 blocks
==12801==      still reachable: 0 bytes in 0 blocks
==12801==            suppressed: 72,704 bytes in 1 blocks
==12801==
==12801== For counts of detected and suppressed errors, rerun with: -v
==12801== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
poohbear@myth30:~/lab2$

```

Not surprisingly, things don't go so well with the buggy version:

```
// reports of still reachable memory resulting from premature exit

poohbear@myth30:~/lab2$ poohbear@myth30:~/lab2$ printf "1 2 3 4" | ./
buggy-exargs factor

poohbear@myth30:~/lab2$ printf "1 2 3 4" | valgrind ./buggy-exargs factor
==14058== Memcheck, a memory error detector
==14058== Copyright (C) 2002-2013, and GNU GPL'd, by Julian Seward et al.
==14058== Using Valgrind-3.10.1 and LibVEX; rerun with -h for copyright info
==14058== Command: ./buggy-exargs factor
==14058==
==14059== Invalid write of size 8
==14059==    at 0x4C2F793: memcpy@@GLIBC_2.14 (in /usr/lib/valgrind/
vgpreload_memcheck-amd64-linux.so)
==14059==    by 0x4013A7: main (buggy-exargs.cc:48)
==14059== Address 0x0 is not stack'd, malloc'd or (recently) free'd
==14059==
==14059==
==14059== Process terminating with default action of signal 11 (SIGSEGV)
==14059== Access not within mapped region at address 0x0
==14059==    at 0x4C2F793: memcpy@@GLIBC_2.14 (in /usr/lib/valgrind/
vgpreload_memcheck-amd64-linux.so)
==14059==    by 0x4013A7: main (buggy-exargs.cc:48)

// reports of still reachable memory resulting from premature exit

poohbear@myth30:~/lab2$
```

`valgrind` tells us that `memcpy` is trying to write memory to an invalid address (`0x0` is as `NULL` an address as they come). Notice that the second process of the second pipeline is set up so that `valgrind` runs `./buggy-exargs` runs `factor`. Neat!

We'll play dumb and pretend we don't know why it's segfaulting, even if `valgrind` does tell us that line 48 of `buggy-exargs.cc` seems to be the culprit. If we want to use `gdb` to set a breakpoint in `buggy-exargs.cc`, line 48, it's totally possible. You just do so like this:


```

poohbear@myth30:~/lab2$ gdb --args ./buggy-exargs factor
// startup preamble omitted for brevity
Reading symbols from ./buggy-exargs...done.
(gdb) list 48
43     vector<string> tokens;
44     pullAllTokens(cin, tokens);
45     pid_t pid = fork();
46     if (pid == 0) {
47         char **exargsv = NULL;
48         memcpy(exargsv, argv + 1, (argc - 1) * sizeof(char *));
49         transform(tokens.cbegin(), tokens.cend(), exargsv + argc - 1,
50                 [](const string& str) { return const_cast<char
*>(str.c_str()); });
51         exargsv[argc + tokens.size() - 1] = NULL;
52         execvp(exargsv[0], exargsv);
(gdb) b 48
Breakpoint 1 at 0x401381: file buggy-exargs.cc, line 48.
(gdb)

```

However, according to [this](#), by default gdb doesn't track execution of additional processes spawned off by the primary. If we were to type run right now, gdb would proceed and fully circumvent the block (or rather, the child process continuation) and miss our breakpoint entirely. Sadness.

Let's confirm by having you execute the following:

- type `run`, advancing the gdb trace of the parent to block on some `getline` call within `pullAllTokens`
- manually type something like "1 2 3 4" (without the double quote) to feed the process's standard input, and then hit ctrl-D to close standard input down
- note that the primary process ends

Here are some questions for you:

- Why is the primary process permitted to run to completion? What is its return value (i.e. its returned status code)?
 - *The primary process runs `exargs`, and `exargs` waits for the seg faulting child to crash*

(and it certainly does). The primary `exargs` returns 1, because that's what it's implemented to do when anything goes wrong with the child.

- Why do we need to manually type in "1 2 3 4"? Why can't we just do what we've done prior and go with something like `printf "1 2 3 4" | gdb --args ./buggy-exargs factor` instead? Why were we able to do this with `valgrind`?
 - *`valgrind` and `strace` ignore standard input (i.e. file descriptor 0 aka `STDIN_FILENO`), whereas `gdb`, which needs to be driven by commands we type in at its prompt, does not. If we try to pipe the output of `printf "1 2 3 4"` to `gdb`, `gdb` accepts 1, 2, 3, and 4 as commands like `run`, `b 48`, and so forth. And once they've been scavenged by `gdb`, they're not available to be read in by `factor`.*
- What `gdb` command can you type in before you type `run` to trace execution of the child process instead of the parent after the `fork`?
 - `set follow-fork-mode yes`
 - *Optionally, set `detach-on-fork` off. By default, `gdb` only pays attention to either parent or child, but not both. If you want `gdb` to track both, you need to turn off process detachment. Reread [this](#) if you want to know more.*

Relaunch `gdb` as you did before, but this time configure it to follow the child process instead of the parent. Confirm that you break on line 48 as intended, and further confirm that `exargsv` is `NULL` as if you've never noticed it before and you've found your bug.